


Apache Spark Fundamentals

OUTLINE

- ① Introduction
- ② Spark Application
- ③ Resilient Distributed Datasets (RDD)
- ④ Implementing of MapReduce Algorithms in Spark
- ⑤ DataFrames and Datasets

Introduction

- Spark is an open-source **framework** for developing and running applications on clusters. (Devised at *UC Berkeley* in 2009 and later donated to *Apache Software Foundation*).
- Spark provides
 - **Unified computing engine** (Spark Core)
 - **Programming interface** usable with **Scala**, **Java**, **Python**, **R**
 - **APIs for data analysis**: **Spark SQL** (structured data) **MLlib** (machine learning), **GraphX** (graph analytics), **Spark Streaming** (streaming analytics). Spark is written in Scala.



- Spark runs on the **Java Virtual Machine (JVM)**


Introduction

- Spark does not come with a storage system (unlike Hadoop) but can run on the *Hadoop Distributed File System (HDFS)* as well as on other systems (e.g., Amazon S3, Cassandra, HIVE, Relational DBMS).
- Spark's features:
 - **Fault tolerance.**
 - **In-memory caching**, which enables **efficient execution of multiround algorithms**, with a substantial performance improvement w.r.t. Hadoop.
- Spark can run:
 - On a single machine in the so called **local mode**.
 - On a cluster managed by a **cluster manager** such as Spark's Standalone, YARN, Mesos.

In the homeworks, you will have the opportunity to use Spark on your machine as well as on a cluster through CloudVeneto (using YARN).

Why do we need a framework like Spark?

Spark Application



Spark Application

- **Driver (a.k.a. master)**: it is the heart of the application, and
 - **Creates the Spark Context**, an object which can be regarded as a channel to access all Spark functionalities.
Obs.: Spark 2.0.0 introduced a Spark Session object which encapsulates the Spark Context and provides more functionalities.
 - **Distributes tasks to the executors.**
 - **Monitors the status of the execution.**

The driver runs Java/Python/Scala/R code through the Spark's API.


- **Executors (a.k.a. workers)**: execute the tasks assigned (**through Scala code**) by the driver, and report their status to the driver.

Spark Application

- **Cluster manager**: when the application is run on a distributed platform, the cluster manager controls the physical machines (i.e., compute nodes) and allocates resources to applications.


EXECUTION: when the application is run on a single machine (i.e., **local mode**), **driver and executors run** on that machine **as threads**. If instead the application is run on a distributed platform (i.e., **cluster mode**) **driver and executors can run on different machines**

MAP/REDUCE TASKS: if the application implements a MapReduce computation, a task assigned to an executor typically comprises running a map function on several key-value pairs or a reduce function on several key-listOfValues pair.



Resilient Distributed Dataset (RDD)

- **Fundamental abstraction** in Spark. An RDD is a collection of elements of the same type, partitioned and distributed across several machines (if available).



- An RDD provides an interface based on **coarse-grained transformations**.
- RDDs ensure **fault-tolerance**.

RDD: Main characteristics

- RDDs are created
 - from data in stable storage (e.g., HDFS), or
 - from other RDDs, through transformations.
- RDDs are immutable (i.e., read-only).
- RDDs are materialized only when needed (lazy evaluation).
- Spark maintains the *lineage* of each RDD, namely the sequence of transformations that generate it, which enable to materialize it or reconstruct it after a failure, starting from data in stable storage.

RDD: Partitioning

- Key ingredient for efficiency: each **RDD is broken into** chunks called **partitions** which are distributed among the available machines.
- A program can specify the **number of partitions for each RDD** (if not, Spark will choose one).
- **Partitions are created by default** (using a **HashPartitioner**, based on objects' hash codes) **or through a custom partitioner**.
- **Typical number of partitions**: is 2x/3x the number of cores, which helps *balancing the work*.
- **Partitioning is exploited to enhance performance**:
 - Spark creates map tasks so to make **each executor apply the map function on data from a locally stored partition** (if possible).
 - To implement algorithms that require partitioning, **the programmer can explicitly access RDD partitions** (e.g., *mapPartition* method).
 - **RDD partitions are exploited implicitly by some ready-made Spark aggregation primitives** (e.g., *reduceByKey* method).

RDD: Operations

The following types of operations can be performed on an RDD *A*

- **TRANSFORMATIONS.** A transformation generates a new RDD *B* starting from the data in *A*. We distinguish between:
 - **Narrow transformation.** Each partition of *A* contributes to (at most) one partition of *B*, which is stored in the same machine. No shuffling of data across machines is needed (\Rightarrow maximum parallelism). E.g., *map* method.
 - **Wide transformation.** Each partition of *A* may contribute to many partitions of *B*. Hence, shuffling of data across machines may be required. E.g., *groupByKey* method.

RDD: Operations

- **ACTIONS.** An action is a computation on the elements of *A* which returns a value to the application. E.g., `count()` method.


LAZY EVALUATION: RDD *A* is **materialized** only when an action is performed. *This must be carefully taken into account when measuring running times*

- **Persistence.** Methods like `cache()` or `persist()` will save the RDD data in memory after the subsequent action.
 - **`cache()`:** data are stored as in RAM (if they fit). Data that do not fit are recomputed when needed.
 - **`persist()`:** data are stored as instructed. There are several options, e.g., `MEMORY_AND_DISK` (RAM+Disk).

RDD: Operations


Narrow transformation

- Input and output stays in same partition
- No data movement is needed



Wide transformation

- Input from other partitions are required
- Data shuffling is needed before processing



Implementing MapReduce algorithms in Spark

The **homeworks** will provide you **first-hand experience** on how to **implement MapReduce algorithms using Spark**.

Summarized here are **key points to keep in mind** about this issue.

- Spark **enables the implementation of MapReduce algorithms** but offers a **much richer set of processing methods**.
- Spark (and MapReduce) exploits a **core idea of functional programming**: *functions can be argument to other functions*.
- A **program which implements a MapReduce algorithm** typically comprises
 - **Spark configuration** steps (e.g., creation of the context);
 - Definitions of **global variables and data structures** (typically, in the driver's memory);
 - **Reading of input data** into one or more RDDs;
 - **Code for the various rounds**.

Implementing a MapReduce round in Spark

Implementing a MapReduce round in Spark

- **Map Phase:** on the input RDD *X* invoke one of map methods offered by Spark (*narrow transformation*) passing the desired map function as argument.
- **Reduce Phase:** on the RDD *X'* resulting from the Map Phase:
 - invoke one of grouping methods offered by Spark (*wide transformation*) to group the key-value pairs into *key-ListOfValues* pairs;
 - invoke one of map methods offered by Spark to apply the desired *reduce function* to each *key-ListOfValues* pair.

Alternatively, one can use one of the ready-made reduce primitives offered by Spark (e.g., `reduceByKey`).

DataFrames and Datasets

- **RDDs** represent the **most basic data model** in Spark: it is **low-level** and **schema-less**.
- On top of RDDs, the **Spark SQL module** provides APIs to operate on the following **structured data** (similar to tables in Relational DB):
 - **DataFrame**: distributed collection of data organized into **named columns**
 - **Dataset**: extension of DataFrame (available only in Java and Scala) with type-safe, object-oriented programming interface.

In the course we use RDDs only!

Summary

- Spark features
- Spark Application: driver process, executor processors, cluster manager.
- Resilient Distributed Dataset (RDD)
 - Main characteristics
 - Partitioning
 - Operations: transformations, actions, persistence.
- Implementing MapReduce algorithms in Spark.
- DataFrames and Datasets.

References

- AS-1 Spark's Web Site: `spark.apache.org`
- AS-2 Spark's RDD Programming guide:
`spark.apache.org/docs/latest/rdd-programming-guide.html`
- CZ18 A Gentle Introduction to Apache Spark. From B. Chambers, M. Zaharia. *Spark: The Definite Guide*, Databricks 2018.
- Z+12 M. Zaharia et al. Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. NSDI 2012: 15-28.