



# CONCURRENT AND REAL TIME PROGRAMMING

[INQ0091623] AA 2021-22

## Lab 3

Working with devices 2

### Video 4 Linux Application Example



Gabriele Manduchi <[gabriele.manduchi@unipd.it](mailto:gabriele.manduchi@unipd.it)>

Andrea Rigoni Garola <[andrea.rigonigarola@unipd.it](mailto:andrea.rigonigarola@unipd.it)>

# Tasks

---


- A Linux application to acquire WebCam images
- Process image to find edges
- We shall use this example to summarize several facts about the Operating System and as a replay for some basic techniques in C programming



# I/O abstraction: A Unified interface for camera devices

( Wikipedia )

**Video4Linux** (V4L for short) is a collection of **device drivers** and an **API** for supporting realtime video capture on Linux systems. It supports many USB webcams, TV tuners, and related devices, standardizing their output, so programmers can easily add video support to applications.



# I/O abstraction: A Unified interface for camera devices

---

- V4L2 (Video for Linux Two), defines a set of ioctl operations and associated data structures.
  - They are general enough to be adapted for all the available camera devices of common usage.
- An important feature is the availability of **query operations** for discovering the supported functionality of the device.
  - To adapt the great variety of camera devices
- V4L2 provides a way to handle a stream of video buffers to pass acquired frames in userspace application.

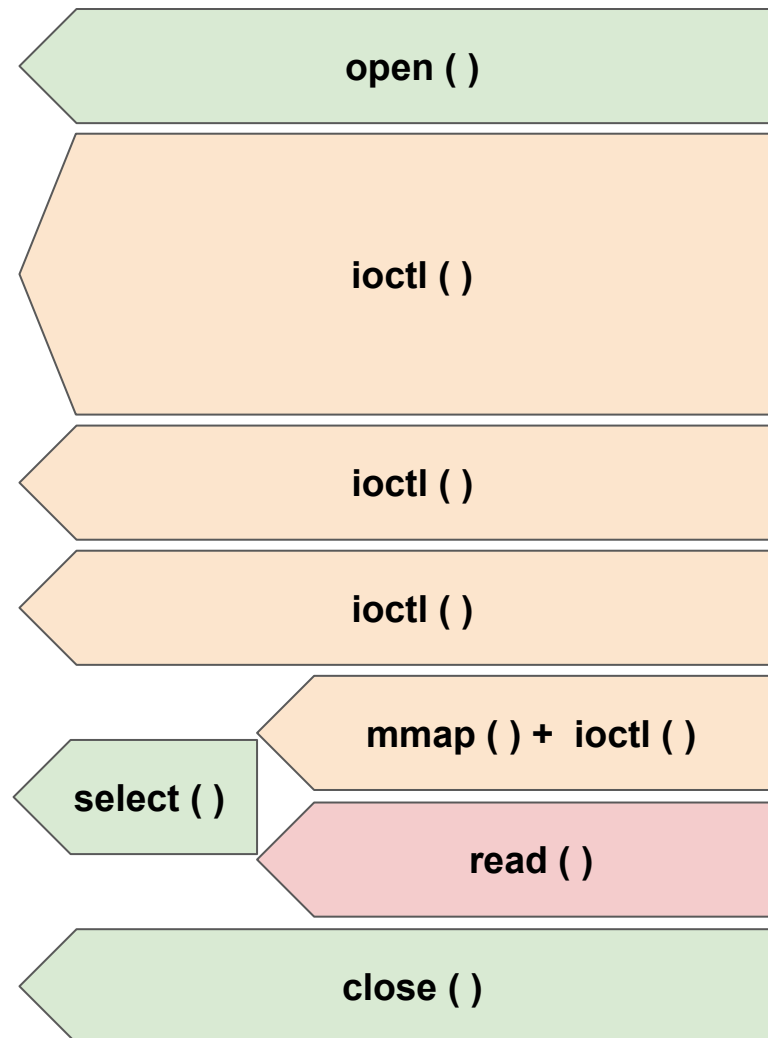
<https://www.linuxtv.org/downloads/v4l-dvb-apis-new/userspace-api/v4l/v4l2.html>

# Common V4L2 API elements

Programming a V4L2 device consists of these steps:

- **Open the device**
- **Change device properties:**
  - selecting a video and audio input
  - video standard
  - picture brightness
  - ...
- **Negotiate a data format**
- **Negotiate an input/output method**
- **Write the actual input/output LOOP**
- **Close the device**

## FILE OPERATIONS



# Opening and Closing Devices

V4L2 drivers are implemented as kernel modules, loaded manually by the system administrator or automatically when a device is first discovered.

The driver modules plug into the videodev kernel module.

Each driver registers one or more device nodes with **major number 81**.  
Minor numbers are allocated dynamically.

```
$ cd /dev/
$ ls -la video*
crw-rw----+ 1 root video 81, 0 Oct 14 04:03 video0
crw-rw----+ 1 root video 81, 1 Oct 14 04:03 video1
```

Default device node name	Usage
/dev/videoX	Video and metadata for capture/output devices
/dev/vbiX	Vertical blank data (i.e. closed captions, teletext)
/dev/radioX	Radio tuners and modulators
/dev/swradioX	Software Defined Radio tuners and modulators
/dev/v4l-touchX	Touch sensors
/dev/v4l-subdevX	Video sub-devices (used by sensors and other components of the hardware peripheral) <sup>1</sup>

# Opening and Closing Devices

---

```
$ v4l2-ctl --list-devices
```

```
UVC Camera (046d:081b)
(usb-0000:00:1a.0-1.3):
    /dev/video0
    /dev/video1
```

```
fd = open("/dev/video0", O_RDWR);

...

close(fd);
```

Which camera is attached to a specific dev file ??

```
$ cd dev/
$ tree v4l
```

```
v4l
|-- by-id
| |-- usb-046d_081b_1F59AD20-video-index0 -> ../../video0
| `-- usb-046d_081b_1F59AD20-video-index1 -> ../../video1
`-- by-path
    |-- pci-0000:00:1a.0-usb-0:1.3:1.0-video-index0 -> ../../video0
    `-- pci-0000:00:1a.0-usb-0:1.3:1.0-video-index1 -> ../../video1
```

# Querying Capabilities

---

Because V4L2 covers a wide variety of devices not all aspects of the API are equally applicable to all types of devices. Furthermore devices of the same type have different capabilities and this specification permits the omission of a few complicated and less important parts of the API.

The [`ioctl VIDIOC\_QUERYCAP`](#) ioctl is available to check if the kernel device is compatible with this specification, and to query the [functions](#) and [I/O methods](#) supported by the device.

V4L2_CAP_READWRITE	0x01000000	The device supports the <b>read()</b> and/or <b>write()</b> I/O methods.
V4L2_CAP_ASYNCIO	0x02000000	The device supports the asynchronous I/O methods.
V4L2_CAP_STREAMING	0x04000000	The device supports the <b>streaming</b> I/O method.

For example in our application we might be interested in checking for streaming input

```
/* Step 2: Check streaming capability */
status = ioctl(fd, VIDIOC_QUERYCAP, &cap);
CHECK_IOCTL_STATUS("Error querying capability")
if(!(cap.capabilities & V4L2_CAP_STREAMING))
{
    printf("Streaming NOT supported\n");
    exit(EXIT_FAILURE);
}
```



## v4l2-ctl - device info

---

Also the **v4l2-ctl** command is able to query many informations from the device:

```
$ v4l2-ctl --info --device /dev/video0
```

```
Driver Info:
```

```
Driver name      : uvcvideo  
Card type       : UVC Camera (046d:081b)  
Bus info        : usb-0000:00:1a.0-1.3  
Driver version  : 5.10.14  
Capabilities    : 0x84a00001
```

```
    Video Capture
```

```
    Metadata Capture
```

```
    Streaming
```

```
    Extended Pix Format
```

```
    Device Capabilities
```

```
Device Caps     : 0x04200001
```

```
    Video Capture
```

```
    Streaming
```

```
    Extended Pix Format
```

# STEP 3 - Query Formats

The word **format** in the V4L context is used to identify how the image is represented in memory

The V4L2 API was primarily designed for devices exchanging **image data** with applications. The struct [v4l2\\_pix\\_format](#) and struct [v4l2\\_pix\\_format\\_mplane](#) structures define the format and layout of an image in memory.

Image formats are negotiated with the [VIDIOC\\_S\\_FMT](#) ioctl.

```
/* Step 3: Check supported formats */
yuyvFound = FALSE;
for(idx = 0; idx < MAX_FORMAT; idx++)
{
    fmt.index = idx;
    fmt.type = V4L2_BUF_TYPE_VIDEO_CAPTURE;
    status = ioctl(fd, VIDIOC_ENUM_FMT, &fmt);
    if(status != 0) break;
    if(fmt.pixelformat == V4L2_PIX_FMT_YUYV)
    {
        yuyvFound = TRUE;
        break;
    }
}
if(!yuyvFound)
{
    printf("YUYV format not supported\n");
    exit(EXIT_FAILURE);
}
```

Query formats using **v4l2-ctl**

```
$ v4l2-ctl --list-formats --device /dev/video0
```

```
ioctl: VIDIOC_ENUM_FMT
Type: Video Capture
```

```
[0]: 'YUYV' (YUYV 4:2:2)
[1]: 'MJPG' (Motion-JPEG, compressed)
```

# Standard image formats

---


To exchange images between the driver and the application, it is necessary to have standard image data formats which both sides will interpret the same way.

V4L2 includes several standard formats already defined. In V4L2 each format has an identifier which looks like **PIX\_FMT\_XXX**, defined in the [videodev2.h](#) header file.

**NOTE:** Custom driver-specific formats are possible. In that case the application may depend on a codec to convert images to one of the standard formats when needed. For example, a device may support a proprietary compressed format. Applications can still capture and save the data in the compressed format, saving much disk space, and later use a codec to convert the images to the X Windows screen format when the video is to be displayed.

The V4L2 standard formats are mainly uncompressed formats.


- The pixels are always arranged in memory **from left to right**, and **from top to bottom**.
- The first byte of data in the image buffer is always for the leftmost pixel of the topmost row.



# RGB

---

These formats encode each pixel as a triplet of RGB values. They are packed formats, meaning that the RGB values for one pixel are stored consecutively in memory and each pixel consumes an integer number of bytes. When the number of bits required to store a pixel is not aligned to a byte boundary, the data is padded with additional bits to fill the remaining byte.

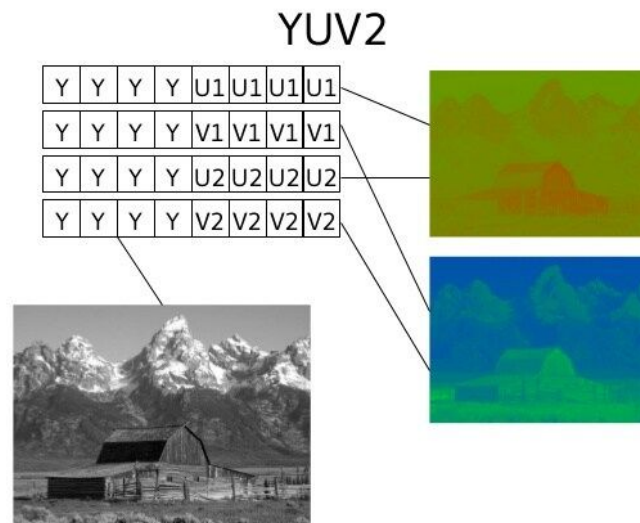


# YUV

YUV is the format native to TV broadcast and composite video signals. It separates the **luminance** brightness information (Y) from the **chrominance** color information (U and V or Cb and Cr). The color information consists of red and blue color difference signals, this way the green component can be reconstructed by subtracting from the brightness component.

$$U = B' - Y' \text{ (blue - luma)}$$
$$V = R' - Y' \text{ (red - luma)}$$

**NOTE:** YUV was chosen because early television would only transmit brightness information. To add color in a way compatible with existing receivers a new signal carrier was added to transmit the color difference signals.



# STEP 4-5-6 - Request Format

CODING EXAMPLE CODING EXAMPLE CODING EXAMPLE

```
/* Step 4: Read current format definition */
memset(&format, 0, sizeof(format));
format.type = V4L2_BUF_TYPE_VIDEO_CAPTURE;
status = ioctl(fd, VIDIOC_G_FMT, &format);
CHECK_IOCTL_STATUS("Error Querying Format")


/* Step 5: Set format fields to desired values: YUYV coding,
480 lines, 640 pixels per line */
format.fmt.pix.width = 640;
format.fmt.pix.height = 480;
format.fmt.pix.pixelformat = V4L2_PIX_FMT_YUYV;

/* Step 6: Write desired format and check actual image size */
status = ioctl(fd, VIDIOC_S_FMT, &format);
CHECK_IOCTL_STATUS("Error Setting Format");
width = format.fmt.pix.width;           //Image Width
height = format.fmt.pix.height;        //Image Height
//Total image size in bytes
imageSize = (unsigned int)format.fmt.pix.sizeimage;
```

# Streaming buffers

Streaming is an I/O method where **only pointers to buffers are exchanged between application and driver**, the **data itself is not actually copied**.

Memory mapping is primarily intended to map buffers in device memory into the application's address space. Device memory can be for example the video memory on a graphics card with a video capture add-on.



First of all we need to ask V4L and the driver to allocate some frames with the `ioctl VIDIOC_REQBUFS`.

struct v4l2_requestbuffers		
__u32	count	The number of buffers requested or granted.
__u32	type	Type of the stream or buffers, this is the same as the struct <code>v4l2_format</code> type field. See <code>v4l2_buf_type</code> for valid values.
__u32	memory	Applications set this field to <code>V4L2_MEMORY_MMAP</code> , <code>V4L2_MEMORY_DMABUF</code> or <code>V4L2_MEMORY_USERPTR</code> . See <code>v4l2_memory</code> .
__u32	capabilities	<p>Set by the driver. If 0, then the driver doesn't support capabilities. In that case all you know is that the driver is guaranteed to support <code>V4L2_MEMORY_MMAP</code> and <i>might</i> support other <code>v4l2_memory</code> types. It will not support any other capabilities.</p> <p>If you want to query the capabilities with a minimum of side-effects, then this can be called with <code>count</code> set to 0, <code>memory</code> set to <code>V4L2_MEMORY_MMAP</code> and <code>type</code> set to the buffer type. This will free any previously allocated buffers, so this is typically something that will be done at the start of the application.</p>
__u32	reserved[1]	A place holder for future extensions. Drivers and applications must set the array to zero.

# STEP 7 - Request Buffer

CODING EXAMPLE CODING EXAMPLE

```
/* Step 7: request for allocation of 4 frame buffers by the driver */
reqBuf.count = 4;
reqBuf.type = V4L2_BUF_TYPE_VIDEO_CAPTURE;
reqBuf.memory = V4L2_MEMORY_MMAP;
status = ioctl(fd, VIDIOC_REQBUFS, &reqBuf);
CHECK_IOCTL_STATUS("Error requesting buffers")
/* Check the number of returned buffers. It must be at least 2 */
if(reqBuf.count < 2)
{
    printf("Insufficient buffers\n");
    exit(EXIT_FAILURE);
}
```



# STEP 8 - Map Buffers

---

To get informations about the actual buffer that the device driver is exposing a special structure is exchanged by the **ioctl** [VIDIOC\\_QUERYBUF](#)

```
struct v4l2_buffer {  
    __u32      index;  
    __u32      type;  
    __u32      bytesused;  
    __u32      flags;  
    __u32      field;  
    struct timeval    timestamp;  
    struct v4l2_timecode    timecode;  
    __u32      sequence;  
};
```

```
/* memory location */  
__u32      memory;  
union {  
    __u32      offset;  
    unsigned long    userptr;  
    struct v4l2_plane    *planes;  
    __s32      fd;  
} m;  
__u32      length;  
__u32      reserved2;  
union {  
    __s32      request_fd;  
    __u32      reserved;  
};  
};
```

# STEP 8 - mmap buffers

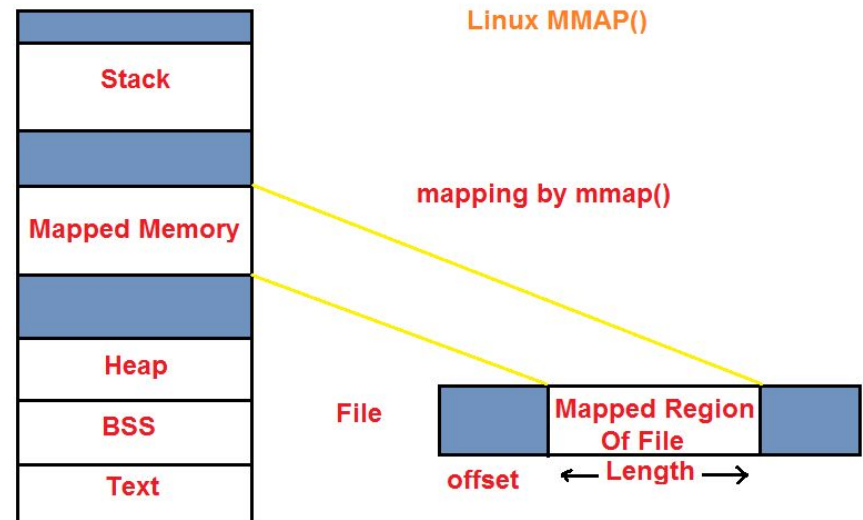
CODING EXAMPLE CODING EXAMPLE

```
struct v4l2_buffer buf;           //Buffer setup structure
buffers = calloc(reqBuf.count, sizeof(bufferDsc));

for(idx = 0; idx < reqBuf.count; idx++)
{
    buf.type = V4L2_BUF_TYPE_VIDEO_CAPTURE;
    buf.memory = V4L2_MEMORY_MMAP;
    buf.index = idx;



    /* Get the start address in the driver space of buffer idx */
    status = ioctl(fd, VIDIOC_QUERYBUF, &buf);
    CHECK_IOCTL_STATUS("Error querying buffers")

    /* Prepare the buffer descriptor with the address in user space
    returned by mmap() */
    buffers[idx].length = buf.length;
    buffers[idx].start = mmap(NULL, buf.length,
        PROT_READ | PROT_WRITE, MAP_SHARED,
        fd, buf.m.offset);
    if(buffers[idx].start == MAP_FAILED)
    {
        perror("Error mapping memory");
        exit(EXIT_FAILURE);
    }
}
```



# Virtual Memory ( REPLAY )

- Provides flexible mapping between the process address space and the physical address space
- Allows different processes run the same program
- May represent an overhead in real-time applications because the context switch required updating Page Table



# Double buffering

---

- In order of avoiding losing frames it is necessary to guarantee that the frame is read before the next one is acquired
- Not feasible in practice, especially for non real-time systems
  - In practice many real-time systems may experience occasional delays
- Using Double buffering , the device parks incoming frames in buffers which are then read by the program
  - In this way, frames are not lost
- Double buffering programming may face the address space problem

# Steps in handling double buffering for camera acquisition

---

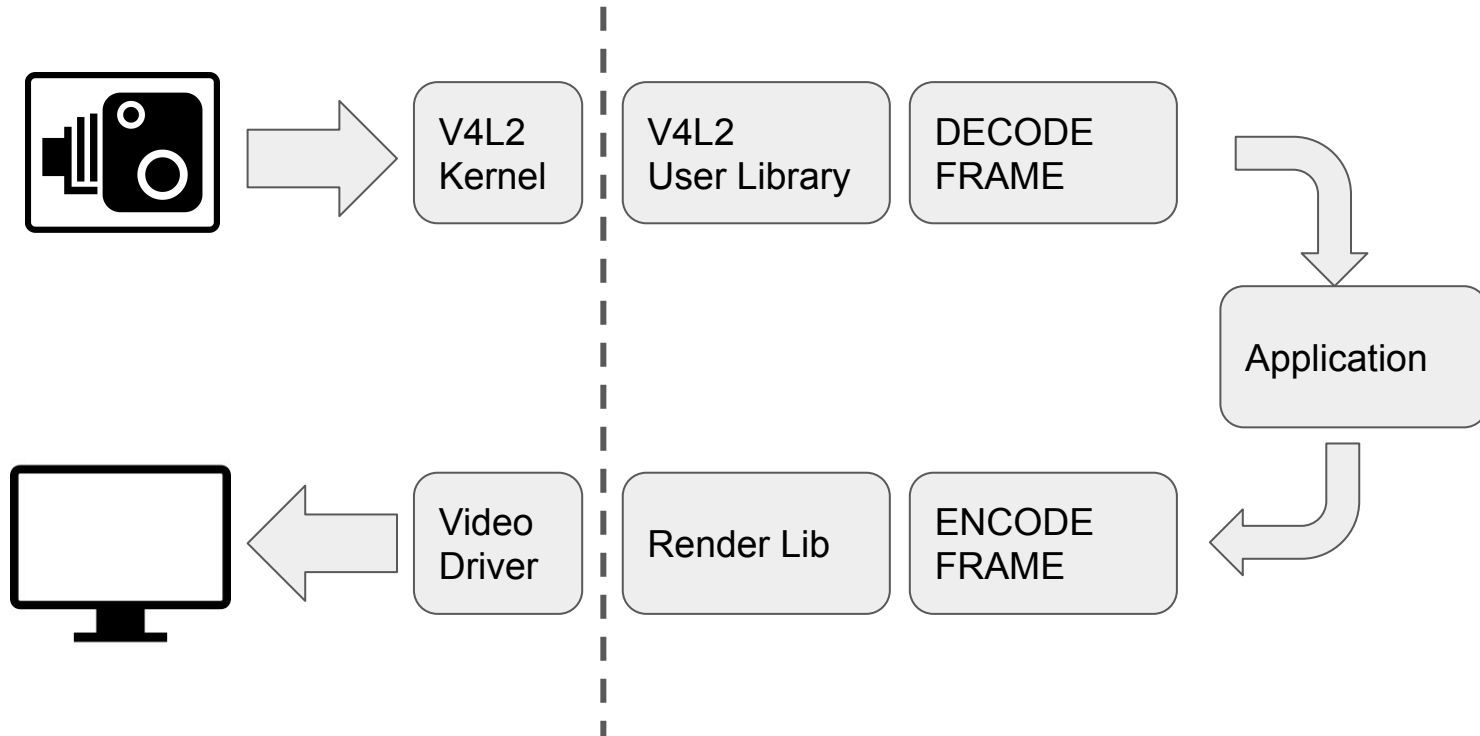
1. Request the device to allocate (in kernel space) a given number of buffers. The buffer offset is returned in the structure associated with `ioctl()` call
2. Map the buffers into process address space using `mmap()`.
3. request the driver to enqueue all the buffers in a circular list
4. Start frame acquisition
5. Wait for buffer ready using `select()`
6. Dequeue the current buffer
7. Use it
8. Enqueue the buffer again

V4L\_example

# Stream of data in a typical video application

---

Look at lab3 files



# Processing image example: Edge detection

---

- Carried out by calculating the gradient  $L_x$  and  $L_y$  over  $X$  and  $Y$ , respectively and selecting the points for which the overall gradient approximated as  $\text{abs}(L_x) + \text{abs}(L_y)$  is above a given threshold
- $X$  and  $Y$  gradient are computed for every pixel by considering a 3x3 matrix (Sobel Matrix)
- A straight implementation uses pointer and memory allocated matrixes to carry out computation



# Sobel Operator

---

The Sobel operator, sometimes called the Sobel–Feldman operator or Sobel filter, is used in image processing and computer vision, particularly within edge detection algorithms where it creates an image emphasising edges. It is named after Irwin Sobel and Gary Feldman, colleagues at the Stanford Artificial Intelligence Laboratory (SAIL).

If we define  $A$  as the source image, and  $G_x$  and  $G_y$  are two images which at each point contain the horizontal and vertical derivative approximations respectively

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

Where  $*$  is the discrete convolution operator

$$g(x, y) = \omega * f(x, y) = \sum_{dx=-a}^a \sum_{dy=-b}^b \omega(dx, dy) f(x + dx, y + dy),$$

$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$