

# Machine Learning

## Linear Models

Fabio Vandin

October 20<sup>th</sup>, 2023


# Linear Classification

$\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{-1, 1\}$ , 0-1 loss

Hypothesis class = *halfspaces*

$$HS_d = \text{sign} \circ L_d = \{\mathbf{x} \rightarrow \text{sign}(h_{\mathbf{w},b}(\mathbf{x})) : h_{\mathbf{w},b} \in L_d\}$$

**Example:**  $\mathcal{X} = \mathbb{R}^2$



# Finding a Good Hypothesis

Linear classification with hypothesis set  $\mathcal{H}$  = halfspaces.

How do we find a good hypothesis? *Good = small generalization error*

# Finding a Good Hypothesis

Linear classification with hypothesis set  $\mathcal{H}$  = halfspaces.

How do we find a good hypothesis?

Good = minimizes the training error (ERM)

$|\mathcal{H}| = +\infty$   
 $\Rightarrow$  "brute force" approach  
does not work

# Finding a Good Hypothesis

Linear classification with hypothesis set  $\mathcal{H}$  = halfspaces.

How do we find a good hypothesis?

Good = minimizes the training error (ERM)

⇒ Perceptron Algorithm (Rosenblatt, 1958)

Training set:  $S = \{(\vec{x}_i, y_i) : 1 \leq i \leq m\}$

**Note:**

if  $y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0$  for all  $i = 1, \dots, m$

hypothesis

$$h_{\vec{w}}(\vec{x}_i) = y_i$$

(prediction made  
by  $h_{\vec{w}}$  on  $\vec{x}_i$   
is "correct", i.e., equal to  $y_i$ )

$h_{\vec{w}}$

# Finding a Good Hypothesis

Linear classification with hypothesis set  $\mathcal{H}$  = halfspaces.

How do we find a good hypothesis?

Good = minimizes the training error (ERM)

$\Rightarrow$  Perceptron Algorithm (Rosenblatt, 1958)

**Note:**

if  $y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0$  for all  $i = 1, \dots, m \Rightarrow$  all points are classified correctly by model  $\mathbf{w} \Rightarrow$  *realizability assumption* for training set

**Linearly separable data:** there exists  $\mathbf{w}$  such that:  $y_i \langle \mathbf{w}, \mathbf{x}_i \rangle > 0$

$$\forall i = 1, \dots, m$$

# Perceptron

**Input:** training set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$

**initialize**  $\mathbf{w}^{(1)} = (0, \dots, 0)$ ;

**for**  $t = 1, 2, \dots$  **do**


**if**  $\exists i$  s.t.  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle \leq 0$  **then**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + y_i \mathbf{x}_i$ ;

**else return**  $\mathbf{w}^{(t)}$ ;

$\mathbf{x}_i$  is misclassified by  $\mathbf{w}^{(t)}$

$\mathbf{w}^{(t)}$  correctly classifies all points in the training set

$$\mathcal{X} = \mathbb{R}^2$$



# Perceptron

**Input:** training set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$


**initialize**  $\mathbf{w}^{(1)} = (0, \dots, 0)$ ;

**for**  $t = 1, 2, \dots$  **do**

**if**  $\exists i$  s.t.  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle \leq 0$  **then**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + y_i \mathbf{x}_i$ ;

**else return**  $\mathbf{w}^{(t)}$ ;

Interpretation of update:



Note that:

$$\begin{aligned} y_i \langle \mathbf{w}^{(t+1)}, \mathbf{x}_i \rangle &= y_i \langle \mathbf{w}^{(t)} + y_i \mathbf{x}_i, \mathbf{x}_i \rangle \\ &= y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle + \underbrace{||\mathbf{x}_i||^2}_{>0} \end{aligned}$$

$\Rightarrow$  update guides  $\mathbf{w}$  to be “more correct” on  $(\mathbf{x}_i, y_i)$ .



# Perceptron

**Input:** training set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$


**initialize**  $\mathbf{w}^{(1)} = (0, \dots, 0)$ ;

**for**  $t = 1, 2, \dots$  **do**

**if**  $\exists i$  s.t.  $y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle \leq 0$  **then**  $\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + y_i \mathbf{x}_i$ ;

**else return**  $\mathbf{w}^{(t)}$ ;

Interpretation of update:



Note that:

$$\begin{aligned} y_i \langle \mathbf{w}^{(t+1)}, \mathbf{x}_i \rangle &= y_i \langle \mathbf{w}^{(t)} + y_i \mathbf{x}_i, \mathbf{x}_i \rangle \\ &= y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle + \|\mathbf{x}_i\|^2 \end{aligned}$$

$\Rightarrow$  update guides  $\mathbf{w}$  to be “more correct” on  $(\mathbf{x}_i, y_i)$ .

Termination? Depends on the realizability assumption!

# Perceptron with Linearly Separable Data

If data is linearly separable one can prove that the perceptron terminates.

## Proposition

Assume that  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$  is linearly separable, let:


- $B = \min\{\|\mathbf{w}\| : y_i \langle \mathbf{w}, \mathbf{x}_i \rangle \geq 1 \ \forall i, i = 1, \dots, m, \}$ , and
- $R = \max_i \|\mathbf{x}_i\|$ .

Then the Perceptron algorithm stops after at most  $(RB)^2$  iterations (and when it stops it holds that  $\forall i, i \in \{1, \dots, m\} : y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle > 0$ ).

# Perceptron: Notes

- simple to implement (but some details are not described in the pseudocode...)
- for separable data
  - termination is guaranteed
  - may require a number of iterations that is exponential in  $d$ ...  
 $\Rightarrow$  other approaches (e.g., ILP - Integer Linear Programming) may be better to find ERM solution in such cases
  - potentially multiple solutions, which one is picked depends on starting values

- non separable data?



# Perceptron: Notes

- simple to implement (but some details are not described in the pseudocode...)
- for separable data
  - termination is guaranteed
  - may require a number of iterations that is exponential in  $d$ ...  
⇒ other approaches (e.g., ILP - Integer Linear Programming) may be better to find ERM solution in such cases
  - potentially multiple solutions, which one is picked depends on starting values
- non separable data?
  - run for some time and keep best solution found up to that point (*pocket algorithm*)


# Perceptron: A Modern View

The previous presentation of the Perceptron is the standard one.

However, we can derive the Perceptron in a different way...

Assume you want to solve a:

- binary classification problem:  $\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{-1, 1\}$
- with linear models
- with loss  $\ell(\vec{w}, (\vec{x}, y)) = \max\{0, -y\langle \vec{w}, \vec{x} \rangle\}$ .



# Perceptron: A Modern View

The previous presentation of the Perceptron is the standard one.

However, we can derive the Perceptron in a different way...

Assume you want to solve a:

- binary classification problem:  $\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{-1, 1\}$
- with linear models
- with loss  $\ell(\mathbf{w}, (\mathbf{x}, y)) = \max\{0, -y\langle \mathbf{w}, \mathbf{x} \rangle\}$ .


Approach: ERM  $\Rightarrow$  need to find the model/hypothesis with smallest training error

How? SGD

**Note:** this is a common framework in all of machine learning!

# Gradient Descent (GD)

General approach for *minimizing* a differentiable convex function



# Gradient Descent (GD)

General approach for *minimizing* a differentiable convex function  $f(\mathbf{w})$

Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  be a differentiable function

## Definition

The *gradient*  $\nabla f(\mathbf{w})$  of  $f$  at  $\mathbf{w} = (w_1, \dots, w_d)$  is

$$\nabla f(\mathbf{w}) = \left( \frac{\partial f(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial f(\mathbf{w})}{\partial w_d} \right)$$

**Intuition:** the gradient points in the direction of the greatest rate of increase of  $f$  around  $\mathbf{w}$



Let  $\eta \in \mathbb{R}, \eta > 0$  be a parameter.

GD algorithm:

$\mathbf{w}^{(0)} \leftarrow \mathbf{0};$

**for**  $t \leftarrow 0$  *to*  $T - 1$  **do**

$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta \nabla f(\mathbf{w}^{(t)});$

**return**  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)};$

Let  $\eta \in \mathbb{R}, \eta > 0$  be a parameter.


GD algorithm:

$\mathbf{w}^{(0)} \leftarrow \mathbf{0}$ ;

for  $t \leftarrow 0$  to  $T - 1$  do

$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta \nabla f(\mathbf{w}^{(t)})$ ;

return  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}$ ;



Notes:

- output vector could also be  $\mathbf{w}^{(T)}$  or  $\arg \min_{\mathbf{w}^{(t)} \in \{1, \dots, T\}} f(\mathbf{w}^{(t)})$
- returning  $\bar{\mathbf{w}}$  is useful for nondifferentiable functions (using *subgradients* instead of gradients...) and for stochastic gradient descent...
- $\eta$ : *learning rate*; sometimes a time dependent  $\eta^{(t)}$  is used (e.g., “move” more at the beginning than at the end)

**Note:** there are guarantees on the number of iterations required by GD to return a *good* value of  $\bar{\mathbf{w}}$  under some assumptions on  $f$  (see the book for details)

# Stochastic Gradient Descent (SGD)

**Idea:** instead of using exactly the gradient, we take a (random) vector with *expected value* equal to the gradient direction.

SGD algorithm:

$\mathbf{w}^{(0)} \leftarrow \mathbf{0}$ ; // or  $\vec{w}^{(0)} \leftarrow \text{random vector}$



for  $t \leftarrow 0$  to  $T-1$  do

    choose  $\mathbf{v}_t$  at random from distribution such that  $\mathbf{E}[\mathbf{v}_t | \mathbf{w}^{(t)}] \in \nabla f(\mathbf{w}^{(t)})$ ;

    /\*  $\mathbf{v}_t$  has *expected value* equal to the gradient of  $f(\mathbf{w}^{(t)})$  \*/

$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta \mathbf{v}_t$ ;

return  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}$ ;



SGD iterations

average of  $\mathbf{w}^{(t)}$

**Note:** there are guarantees on the number of iterations required by GD to return a *good, in expectation*, value of  $\bar{\mathbf{w}}$  under some assumptions on  $f$  (see the book for details)

Why should we use SGD instead of GD?

**Question:** when do we use GD in the first place?

**Answer:** for example to find  $\mathbf{w}$  that minimizes  $L_S(\mathbf{w})$

That is: we use GD for  $f(\mathbf{w}) = L_S(\mathbf{w})$

$\Rightarrow \nabla f(\mathbf{w})$  depends on all pairs  $(\mathbf{x}_i, y_i) \in S, i = 1, \dots, m$ : may require long time to compute it!

**What about SGD?**

We need to pick  $\mathbf{v}_t$  such that  $\mathbf{E}[\mathbf{v}_t | \mathbf{w}^{(t)}] \in \nabla f(\mathbf{w}^{(t)})$ : **how?**

Pick a random  $(\mathbf{x}_i, y_i) \in S \Rightarrow$  pick  $\mathbf{v}_t \in \nabla \ell(\mathbf{w}^{(t)}, (\mathbf{x}_i, y_i))$ :

- satisfies the requirement!
- requires much less computation than GD

Analogously we can use SGD for regularized losses, etc.

# Back to Our Linear Classification Problem

- binary classification problem:  $\mathcal{X} = \mathbb{R}^d$ ,  $\mathcal{Y} = \{-1, 1\}$
- with linear models
- with loss  $\ell(\mathbf{w}, (\mathbf{x}, y)) = \max\{0, -y\langle \mathbf{w}, \mathbf{x} \rangle\}$ .

How to find the ERM solution? SGD!

## SGD for Linear Classification

SGD: take  $i$  uniformly at random from  $\{1, \dots, m\}$ .

Let  $(\vec{x}', y')$  be the corresponding point in the training set, and consider the vector  $\nabla \ell(\vec{w}, (\vec{x}', y'))$

Note that GD considers (as gradient of the function to minimize):

$$\nabla L_S(\vec{w}) = \frac{1}{m} \sum_{i=1}^m \nabla \ell(\vec{w}, (\vec{x}_i, y_i))$$

and for SGD we have:  $\frac{1}{m} \forall i \in \{1, \dots, m\}$  (uniform distribution)

$$E[\nabla \ell(\vec{w}, (\vec{x}', y'))] = \sum_{i=1}^m \underbrace{\Pr[(\vec{x}', y') = (\vec{x}_i, y_i)]}_{\frac{1}{m}} \cdot \nabla \ell(\vec{w}, (\vec{x}_i, y_i))$$

$$= \frac{1}{m} \sum_{i=1}^m \nabla \ell(\vec{w}, (\vec{x}_i, y_i))$$

$$= \nabla L_S(\vec{w})$$